

# OS LAB MANUAL

## Experiment 1: Practicing of Basic UNIX Commands

### **1. pwd (Print Working Directory)**

Displays the current directory path.

```
$ pwd
```

```
/home/user
```

---

### **2. ls (List Directory Contents)**

Lists files and directories in the current directory.

```
$ ls
```

```
file1.txt file2.txt folder1
```

- **Options:**

- `ls -l`: Long format (shows details like permissions, size, and date).
  - `ls -a`: Includes hidden files.
- 

### **3. cd (Change Directory)**

Changes to a different directory.

```
$ cd folder1
```

```
$ pwd
```

```
/home/user/folder1
```

- `cd ..`: Moves to the parent directory.
- 

### **4. mkdir (Make Directory)**

Creates a new directory.

```
$ mkdir new_folder
```

```
$ ls
```

new\_folder

---

## **5. touch (Create Empty Files)**

Creates a new, empty file.

```
$ touch new_file.txt
```

```
$ ls
```

```
new_file.txt
```

---

## **6. cat (Concatenate and Display Files)**

Displays the contents of a file.

```
$ cat file1.txt
```

```
Hello, this is a file.
```

---

## **7. cp (Copy Files and Directories)**

Copies files or directories.

```
$ cp file1.txt file2.txt
```

```
$ ls
```

```
file1.txt file2.txt
```

---

## **8. mv (Move or Rename Files)**

Moves or renames files.

```
$ mv file1.txt new_name.txt
```

```
$ ls
```

```
new_name.txt
```

---

## **9. rm (Remove Files or Directories)**

Deletes files or directories.

```
$ rm file2.txt
```

```
$ ls
```

```
new_name.txt
```

- **Options:**

- `rm -r folder_name`: Removes a directory and its contents.
- 

## **10. man (Manual Pages)**

Shows the manual for a command.

```
$ man ls
```

---

## **11. chmod (Change File Permissions)**

Changes file permissions.

```
$ chmod 644 file.txt
```

- Example permissions:

- 644: Read and write for owner, read-only for others.
- 

## **12. whoami**

Displays the current user.

```
$ whoami
```

```
user
```

---

## **13. clear**

Clears the terminal screen.

```
$ clear
```

---

## **14. find**

Searches for files or directories.

```
$ find . -name "file1.txt"
./folder1/file1.txt
```

---

## 15. grep (Search for Text in Files)

Searches for a specific string in a file.

```
$ grep "text" file.txt
```

This is the text you searched for.

---

## 16. df (Disk Free)

Displays disk space usage.

```
$ df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/sda1	50G	20G	30G	40%	/

---

## 17. top

Shows real-time system processes.

```
$ top
```

---

## 18. exit

Logs out of the current shell session.

```
$ exit
```

**Experiment 2: Write programs using the following UNIX operating system calls fork, exec, getpid, exit, wait, close, stat, opendir and readdir**

### 1. fork and getpid

Creates a child process and displays the process IDs of parent and child.

**Program:**

```
#include <stdio.h>
```

```
#include <unistd.h>
```

```
int main() {
    pid_t pid = fork();
    if (pid == 0) {
        // Child process
        printf("Child process ID: %d\n", getpid());
    } else if (pid > 0) {
        // Parent process
        printf("Parent process ID: %d\n", getpid());
    } else {
        perror("Fork failed");
    }
    return 0;
}
```

---

## 2. exec

Executes a new program (ls command) from the current process.

### Program:

```
#include <stdio.h>
#include <unistd.h>

int main() {
    printf("Executing `ls` command:\n");
    execl("/bin/ls", "ls", "-l", NULL);
    perror("Exec failed");
    return 0;
}
```

---

## 3. exit and wait

Demonstrates a parent process waiting for the child to exit.

**Program:**

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main() {
    pid_t pid = fork();
    if (pid == 0) {
        printf("Child process exiting.\n");
        exit(0);
    } else if (pid > 0) {
        int status;
        wait(&status);
        printf("Parent process: Child exited with status %d.\n",
            WEXITSTATUS(status));
    } else {
        perror("Fork failed");
    }
    return 0;
}
```

---

#### **4. close**

Opens and closes a file descriptor.

**Program:**

```
#include <stdio.h>
#include <fcntl.h>
#include <unistd.h>
```

```
int main() {
int fd = open("test.txt", O_CREAT | O_RDWR, 0644);
if (fd < 0) {
perror("Failed to open file");
return 1;
}
printf("File opened with descriptor: %d\n", fd);
close(fd);
printf("File descriptor closed.\n");
return 0;
}
```

---

## 5. stat

Displays information about a file.

### Program:

```
#include <stdio.h>
#include <sys/stat.h>
int main() {
struct stat fileStat;
if (stat("test.txt", &fileStat) < 0) {
perror("Failed to get file stats");
return 1;
}

printf("File size: %ld bytes\n", fileStat.st_size);
printf("Permissions: %o\n", fileStat.st_mode & 0777);
```

```
printf("Last accessed: %ld\n", fileStat.st_atime);  
return 0;  
}
```

---

## 6. opendir and readdir

Lists the contents of a directory.

### Program:

```
#include <stdio.h>  
#include <dirent.h>  
  
int main() {  
    DIR *dir = opendir(".");  
    if (dir == NULL) {  
        perror("Failed to open directory");  
        return 1;  
    }  
  
    struct dirent *entry;  
    while ((entry = readdir(dir)) != NULL) {  
        printf("%s\n", entry->d_name);  
    }  
  
    closedir(dir);  
    return 0;  
}
```

---

1. Save the code to a file, e.g., program.c.



2. Compile using gcc:

```
gcc program.c -o program
```

3. Run the program:

```
./program
```

### **Experiment 3:Simulate UNIX commands like cp, ls, grep, etc.,**

#### **Simulating cp (Copying Files)**

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
// Function to copy the content of the source file to the destination file
```

```
void copy_file(char *source, char *destination) {
```

```
    FILE *src, *dest;
```

```
    char ch;
```

```
    // Open the source file in read mode
```

```
    src = fopen(source, "r");
```

```
    if (src == NULL) {
```

```
        perror("Source file opening failed");
```

```
        return;
```

```
    }
```

```
    // Open the destination file in write mode
```

```
    dest = fopen(destination, "w");
```

```
    if (dest == NULL) {
```

```
        perror("Destination file opening failed");
```

```
        fclose(src);
```

```
        return;
```

```
    }
```

```
    // Copy each character from source file to destination file
    while ((ch = fgetc(src)) != EOF) {
fputc(ch, dest);
    }
```

```
printf("File copied successfully.\n");
```

```
    // Close both the source and destination files
fclose(src);
fclose(dest);
}
```

```
int main() {
    char source[100], destination[100];
```

```
    // Prompt user for source and destination file names
printf("Enter source file name: ");
scanf("%s", source);
printf("Enter destination file name: ");
scanf("%s", destination);
```

```
    // Call the copy_file function to copy the content
copy_file(source, destination);
```

```
    return 0;
}
```

## Simulating ls (Listing Files in a Directory)

To simulate the ls command, we can use the opendir and readdir functions from the <dirent.h>

```
#include <stdio.h>

#include <dirent.h>

// Function to list all files in the given directory
void list_files(char *path) {
    struct dirent *entry;
    DIR *dir = opendir(path);

    if (dir == NULL) {
        perror("Directory opening failed");
        return;
    }

    printf("Files in %s:\n", path);
    while ((entry = readdir(dir)) != NULL) {
        printf("%s\n", entry->d_name);
    }

    closedir(dir);
}

int main() {
    char path[100];

    printf("Enter directory path: ");
```

```
scanf("%s", path); // Added missing semicolon here
```

```
list_files(path); // Added missing semicolon here
```

```
    return 0;  
}
```

### **Simulating grep (Searching for a Pattern in a File)**

This program will search for a specific string (pattern) in a file, similar to how grep works in UNIX.

```
#include <stdio.h>
```

```
#include <string.h>
```

```
// Function to search for a pattern in a file and print matching lines
```

```
void grep_pattern(char *file_name, char *pattern) {
```

```
    FILE *file;
```

```
    char line[256];
```

```
    int line_number = 0;
```

```
    file = fopen(file_name, "r");
```

```
    if (file == NULL) {
```

```
perror("File opening failed");
```

```
        return;
```

```
    }
```

```
    while (fgets(line, sizeof(line), file)) {
```

```
line_number++;
```

```
        if (strstr(line, pattern)) {
```

```
printf("Line %d: %s", line_number, line);
```

```

    }
}

fclose(file);
}

int main() {
    char file_name[100], pattern[100];

    printf("Enter file name: ");
    scanf("%s", file_name);
    printf("Enter pattern to search: ");
    scanf("%s", pattern);

    grep_pattern(file_name, pattern);

    return 0; // Properly terminate the main function
}

```

**Experiment 4: Simulate the following CPU scheduling algorithms a) FCFS  
b) SJF c) Priority d) Round Robin**

### **1. First-Come, First-Served (FCFS):**

FCFS is the simplest scheduling algorithm where the process that arrives first is executed first.

```
#include <stdio.h>
```

```
// Function to implement FCFS scheduling
```

```
void FCFS(int n, int burst_time[]) {
    int wait_time = 0, turnaround_time = 0;
```

```

float avg_wait_time = 0, avg_turnaround_time = 0;

for (int i = 0; i < n; i++) {
wait_time += turnaround_time;
turnaround_time += burst_time[i];
avg_wait_time += wait_time;
avg_turnaround_time += turnaround_time;
}

printf("FCFS Scheduling:\n");
printf("Average Waiting Time: %.2f\n", avg_wait_time / n);
printf("Average Turnaround Time: %.2f\n", avg_turnaround_time / n);
}

int main() {
    int n;

    printf("Enter the number of processes: ");
    scanf("%d", &n);

    int burst_time[n];

    printf("Enter the burst times of the processes:\n");
    for (int i = 0; i < n; i++) {
        printf("Process %d: ", i + 1);
        scanf("%d", &burst_time[i]);
    }

```

```
FCFS(n, burst_time);
```

```
    return 0;  
}
```

## **2. Shortest Job First (SJF):**

SJF schedules processes with the shortest burst time first.

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
// Function to implement Shortest Job First (SJF) scheduling
```

```
void SJF(int n, int burst_time[]) {
```

```
    int wait_time = 0, turnaround_time = 0;
```

```
    float avg_wait_time = 0, avg_turnaround_time = 0;
```

```
    // Sort burst times in ascending order
```

```
    for (int i = 0; i < n - 1; i++) {
```

```
        for (int j = 0; j < n - i - 1; j++) {
```

```
            if (burst_time[j] > burst_time[j + 1]) {
```

```
                int temp = burst_time[j];
```

```
                burst_time[j] = burst_time[j + 1];
```

```
                burst_time[j + 1] = temp;
```

```
            }
```

```
        }
```

```
    }
```

```
    for (int i = 0; i < n; i++) {
```

```
wait_time += turnaround_time;
turnaround_time += burst_time[i];
avg_wait_time += wait_time;
avg_turnaround_time += turnaround_time;
}
```

```
printf("SJF Scheduling:\n");
printf("Average Waiting Time: %.2f\n", avg_wait_time / n);
printf("Average Turnaround Time: %.2f\n", avg_turnaround_time / n);
}
```

```
int main() {
```

```
    int n;
```

```
    printf("Enter the number of processes: ");
```

```
    scanf("%d", &n);
```

```
    int burst_time[n];
```

```
    printf("Enter the burst times of the processes:\n");
```

```
    for (int i = 0; i < n; i++) {
```

```
        printf("Process %d: ", i + 1);
```

```
        scanf("%d", &burst_time[i]); // Fixed missing semicolon
```

```
    }
```

```
    SJF(n, burst_time);
```



```
    return 0; // Properly terminate the main function
}
```

### 3. Priority Scheduling:

Priority scheduling selects the process with the highest priority (lowest numerical value) first.

```
#include <stdio.h>
```

```
// Function to implement Priority Scheduling
```

```
void Priority(int n, int burst_time[], int priority[]) {
    int wait_time = 0, turnaround_time = 0;
    float avg_wait_time = 0, avg_turnaround_time = 0;
    int temp_burst, temp_priority;

    // Sort processes based on priority
    for (int i = 0; i < n - 1; i++) {
        for (int j = 0; j < n - i - 1; j++) {
            if (priority[j] > priority[j + 1]) {
                // Swap burst times
                temp_burst = burst_time[j];
                burst_time[j] = burst_time[j + 1];
                burst_time[j + 1] = temp_burst;

                // Swap priorities
                temp_priority = priority[j];
                priority[j] = priority[j + 1];
                priority[j + 1] = temp_priority;
            }
        }
    }
}
```

```

    }

    // Calculate waiting and turnaround times
    for (int i = 0; i < n; i++) {
wait_time += turnaround_time;
turnaround_time += burst_time[i];
avg_wait_time += wait_time;
avg_turnaround_time += turnaround_time;
    }

printf("Priority Scheduling:\n");
printf("Average Waiting Time: %.2f\n", avg_wait_time / n);
printf("Average Turnaround Time: %.2f\n", avg_turnaround_time / n);
}

int main() {
    int n;

printf("Enter the number of processes: ");
scanf("%d", &n);

    int burst_time[n], priority[n];

printf("Enter the burst times and priorities of the processes:\n");
    for (int i = 0; i < n; i++) {
printf("Process %d:\n", i + 1);
printf("Burst time: ");

```

```
scanf("%d", &burst_time[i]);
printf("Priority: ");
scanf("%d", &priority[i]);
}
```

```
Priority(n, burst_time, priority);
```

```
return 0;
}
```

#### **4. Round Robin (RR):**

Round Robin scheduling allocates a fixed time quantum for each process. If a process doesn't complete within its time slice, it's moved to the back of the queue.

```
#include <stdio.h>

void RoundRobin(int n, int burst_time[], int quantum) {
    int wait_time = 0, turnaround_time = 0;
    int remaining_burst_time[n];
    float avg_wait_time = 0, avg_turnaround_time = 0;
    // Copy burst time to remaining_burst_time
    for (int i = 0; i < n; i++) {
        remaining_burst_time[i] = burst_time[i];
    }
    int time = 0;
    while (1) {
        int completed = 1;
        for (int i = 0; i < n; i++) {
            if (remaining_burst_time[i] > 0) {
                completed = 0;
```

```

if (remaining_burst_time[i] > quantum) {
time += quantum;
remaining_burst_time[i] -= quantum;
} else {
time += remaining_burst_time[i];
wait_time += time - burst_time[i];
remaining_burst_time[i] = 0;
}
}
}
if (completed) break;
}
for (int i = 0; i < n; i++) {
turnaround_time += burst_time[i];
avg_turnaround_time += turnaround_time;
}
avg_wait_time = (float)wait_time / n;
avg_turnaround_time = (float)avg_turnaround_time / n;
printf("Round Robin Scheduling (Quantum = %d):\n", quantum);
printf("Average Waiting Time: %.2f\n", avg_wait_time);
printf("Average Turnaround Time: %.2f\n", avg_turnaround_time);
}

int main() {
int n, quantum;
printf("Enter the number of processes: ");
scanf("%d", &n);
int burst_time[n];

```

```

printf("Enter the burst times of the processes:\n");
for (int i = 0; i < n; i++) {
printf("Process %d: ", i + 1);
scanf("%d", &burst_time[i]);
}
printf("Enter the time quantum: ");
scanf("%d", &quantum);
RoundRobin(n, burst_time, quantum);
return 0;
}

```

### **Compile the programs using gcc:**

```

gcc FCFS.c -o FCFS
gcc SJF.c -o SJF
gcc Priority.c -o Priority
gcc RR.c -o RR

```

### **Run the executable:**

```

./FCFS
./SJF
./Priority
./RR

```

## **Experiment 5:Control the number of ports opened by the operating system with a) Semaphore b) Monitors.**

### **1. Using Semaphore to Control the Number of Ports Opened**

A **semaphore** is a synchronization primitive that can be used to control access to shared resources. The semaphore works by maintaining a counter that represents the number of available resources. If the counter is positive, processes can acquire a resource; if the counter is zero, processes must wait until the resource becomes available.

### **Example of Controlling Ports Using Semaphore**

```
#include <stdio.h>
#include <stdlib.h>
#include <semaphore.h>
#include <pthread.h>
#include <unistd.h> // For sleep function

#define MAX_PORTS 3 // Maximum number of ports that can be opened

sem_t available_ports; // Semaphore to control access to ports

void* open_port(void* id) {
    int thread_id = *(int*)id;
    printf("Thread %d: Attempting to open a port...\n", thread_id);

    // Wait for an available port
    sem_wait(&available_ports);
    printf("Thread %d: Port opened.\n", thread_id);

    // Simulate port usage
    sleep(2);

    // Release the port
    printf("Thread %d: Closing port.\n", thread_id);
    sem_post(&available_ports);

    return NULL;
}
```

```

int main() {
pthread_t threads[5]; // 5 threads, trying to open ports
    int thread_ids[5];

    // Initialize semaphore with MAX_PORTS available ports
sem_init(&available_ports, 0, MAX_PORTS);

    // Create 5 threads, each trying to open a port
    for (int i = 0; i < 5; i++) {
thread_ids[i] = i + 1;
pthread_create(&threads[i], NULL, open_port, &thread_ids[i]);
    }

    // Wait for threads to finish
    for (int i = 0; i < 5; i++) {
pthread_join(threads[i], NULL);
    }

    // Destroy the semaphore
sem_destroy(&available_ports);

    return 0;
}

```

## 2. Using Monitors to Control the Number of Ports Opened

A **monitor** is a higher-level synchronization construct that encapsulates both data and methods for manipulating that data. It ensures that only one thread can execute the critical section of code at a time. In C, we can use mutexes or

condition variables (through POSIX threads or pthread library) to implement monitors.

### **Example of Controlling Ports Using Monitors**

Here, we will implement a monitor-like behavior using mutexes and condition variables to simulate controlling the number of ports.

```
#include <stdio.h>

#include <stdlib.h>

#include <pthread.h>

#include <unistd.h> // For sleep function


#define MAX_PORTS 3 // Maximum number of ports that can be opened


pthread_mutex_t mutex; // Mutex to protect shared resource (ports)
pthread_cond_t cond; // Condition variable for synchronization
int available_ports = MAX_PORTS; // Shared resource (number of available
ports)


void* open_port(void* id) {
    int thread_id = *(int*)id;
    printf("Thread %d: Attempting to open a port...\n", thread_id);

    pthread_mutex_lock(&mutex); // Enter critical section

    // Wait if no ports are available
    while (available_ports == 0) {
        pthread_cond_wait(&cond, &mutex);
    }
```



```

    // Open a port
    available_ports--;

    printf("Thread %d: Port opened. Remaining ports: %d\n", thread_id,
    available_ports);

    pthread_mutex_unlock(&mutex); // Exit critical section


    // Simulate port usage
    sleep(2);


    pthread_mutex_lock(&mutex); // Enter critical section to close port
    available_ports++; // Release the port
    printf("Thread %d: Closing port. Remaining ports: %d\n", thread_id,
    available_ports);


    pthread_cond_signal(&cond); // Signal other threads that a port is available
    pthread_mutex_unlock(&mutex); // Exit critical section


    return NULL;
}


int main() {
    pthread_t threads[5]; // 5 threads, trying to open ports
    int thread_ids[5];


    // Initialize the mutex and condition variable
    pthread_mutex_init(&mutex, NULL);
    pthread_cond_init(&cond, NULL);

```

```

    // Create 5 threads, each trying to open a port
    for (int i = 0; i < 5; i++) {
thread_ids[i] = i + 1;
pthread_create(&threads[i], NULL, open_port, &thread_ids[i]);
    }

    // Wait for threads to finish
    for (int i = 0; i < 5; i++) {
pthread_join(threads[i], NULL);
    }

    // Destroy the mutex and condition variable
pthread_mutex_destroy(&mutex);
pthread_cond_destroy(&cond);

    return 0;
}

```

**Experiment 6: Write a program to illustrate concurrent execution of threads using threads library.**

```

#include <stdio.h>
#include <pthread.h>
#include <unistd.h>

void* print_message(void* thread_id) {
    long tid = (long) thread_id;
    printf("Thread %ld is running\n", tid);
    sleep(2); // Simulate some work by sleeping for 2 seconds
    printf("Thread %ld is done\n", tid);
}

```

```

        return NULL;
    }

int main() {
pthread_t threads[5]; // Array to hold thread identifiers
    int num_threads = 5;

    // Create 5 threads
    for (long i = 0; i < num_threads; i++) {
        int result = pthread_create(&threads[i], NULL, print_message, (void*)i);
        if (result != 0) {
perror("Thread creation failed");
            return -1;
        }
    }

    // Wait for all threads to finish
    for (int i = 0; i < num_threads; i++) {
pthread_join(threads[i], NULL);
    }

    printf("All threads have finished execution.\n");
    return 0;
}

```

**Experiment 7: Write a program to solve producer-consume problem using Semaphores.**

```

#include <stdio.h>

#include <stdlib.h>

```

```

#include <pthread.h>
#include <semaphore.h>
#include <unistd.h>

#define BUFFER_SIZE 5 // Size of the shared buffer
#define PRODUCER_COUNT 3 // Number of producers
#define CONSUMER_COUNT 3 // Number of consumers

int buffer[BUFFER_SIZE]; // Shared buffer
int in = 0, out = 0; // Indices for producer and consumer
sem_t empty, full, mutex; // Semaphores

void* producer(void* arg) {
    for (int i = 0; i < 5; i++) {
        int item = rand() % 100; // Produce a random item
        sem_wait(&empty); // Decrement empty (wait for space)
        sem_wait(&mutex); // Enter critical section
        buffer[in] = item; // Add item to buffer
        printf("Producer %ld produced item: %d at index %d\n", (long)arg, item, in);
        in = (in + 1) % BUFFER_SIZE; // Move to next position
        sem_post(&mutex); // Exit critical section
        sem_post(&full); // Increment full (notify consumer)
        sleep(1); // Simulate time taken to produce
    }
    return NULL;
}

```

```

void* consumer(void* arg) {
    for (int i = 0; i < 5; i++) {
sem_wait(&full); // Decrement full (wait for item)
sem_wait(&mutex); // Enter critical section

        int item = buffer[out]; // Consume an item from the buffer
printf("Consumer %ld consumed item: %d from index %d\n", (long)arg, item,
out);

        out = (out + 1) % BUFFER_SIZE; // Move to next position
sem_post(&mutex); // Exit critical section
sem_post(&empty); // Increment empty (notify producer)
sleep(2); // Simulate time taken to consume
    }
    return NULL;
}

```

```

int main() {
pthread_t producers[PRODUCER_COUNT],
consumers[CONSUMER_COUNT];

    // Initialize semaphores
sem_init(&empty, 0, BUFFER_SIZE); // Initially, the buffer is empty
sem_init(&full, 0, 0);           // Initially, no items are produced
sem_init(&mutex, 0, 1);          // Mutex for mutual exclusion in buffer

    // Create producer threads
    for (long i = 0; i < PRODUCER_COUNT; i++) {
pthread_create(&producers[i], NULL, producer, (void*)i);
    }
}

```

```

// Create consumer threads
for (long i = 0; i < CONSUMER_COUNT; i++) {
pthread_create(&consumers[i], NULL, consumer, (void*)i);
}

// Wait for all producer threads to finish
for (int i = 0; i < PRODUCER_COUNT; i++) {
pthread_join(producers[i], NULL);
}

// Wait for all consumer threads to finish
for (int i = 0; i < CONSUMER_COUNT; i++) {
pthread_join(consumers[i], NULL);
}

// Destroy semaphores
sem_destroy(&empty);
sem_destroy(&full);
sem_destroy(&mutex);

return 0;
}

```

**Experiment 8: Implement the following memory allocation methods for fixed partition a) First fit b) Worst fit c) Best fit**

### **1. First Fit Algorithm**

```

#include <stdio.h>

#define PARTITION_COUNT 5 // Number of fixed partitions

```

```

#define PROCESS_COUNT 4 // Number of processes

// Function to perform First Fit allocation
void firstFit(int partitions[], int partitionCount, int processes[], int
processCount) {

    int allocation[processCount]; // Declare allocation array with correct size

    // Initially, no process is allocated
    for (int i = 0; i<processCount; i++) {
        allocation[i] = -1;
    }

    // Try to allocate memory for each process
    for (int i = 0; i<processCount; i++) {
        for (int j = 0; j <partitionCount; j++) {
            if (partitions[j] >= processes[i]) {
                allocation[i] = j;
                partitions[j] -= processes[i]; // Allocate the partition
                break;
            }
        }
    }

    // Output the result
    printf("\nFirst Fit Allocation:\n");
    for (int i = 0; i<processCount; i++) {
        if (allocation[i] != -1) {
            printf("Process %d allocated to Partition %d\n", i + 1, allocation[i] + 1);
        }
    }
}

```

```

        } else {
printf("Process %d not allocated\n", i + 1);
        }
    }
}

int main() {
    int partitions[PARTITION_COUNT] = {100, 500, 200, 300, 600}; // Partition
    sizes
    int processes[PROCESS_COUNT] = {212, 417, 112, 426}; // Process
    sizes
    firstFit(partitions, PARTITION_COUNT, processes, PROCESS_COUNT);
    return 0;
}

```

## 2. Worst Fit Algorithm

```

#include <stdio.h>

#define PARTITION_COUNT 5 // Number of fixed partitions
#define PROCESS_COUNT 4 // Number of processes

// Function to perform Worst Fit allocation
void worstFit(int partitions[], int partitionCount, int processes[], int
processCount) {
    int allocation[processCount]; // Allocation array to store partition
    assignments

    // Initially, no process is allocated
    for (int i = 0; i < processCount; i++) {
        allocation[i] = -1;
    }
}

```



```

// Try to allocate memory for each process
for (int i = 0; i < processCount; i++) {
    int maxIndex = -1;
    int maxSize = -1;

    // Find the partition with the largest size that can accommodate the process
    for (int j = 0; j < partitionCount; j++) {
        if (partitions[j] >= processes[i] && partitions[j] > maxSize) {
            maxSize = partitions[j];
            maxIndex = j;
        }
    }

    // If a suitable partition is found, allocate it
    if (maxIndex != -1) {
        allocation[i] = maxIndex;
        partitions[maxIndex] -= processes[i]; // Allocate the partition
    }
}

// Output the result
printf("\nWorst Fit Allocation:\n");
for (int i = 0; i < processCount; i++) {
    if (allocation[i] != -1) {
        printf("Process %d allocated to Partition %d\n", i + 1, allocation[i] + 1);
    } else {

```

```

printf("Process %d not allocated\n", i + 1);
    }
}
}

int main() {
    int partitions[PARTITION_COUNT] = {100, 500, 200, 300, 600}; // Partition
    sizes
    int processes[PROCESS_COUNT] = {212, 417, 112, 426}; // Process
    sizes
    worstFit(partitions, PARTITION_COUNT, processes, PROCESS_COUNT);
    return 0;
}

```

### **3.Best Fit Algorithm**

```

#include <stdio.h>

#define PARTITION_COUNT 5 // Number of fixed partitions
#define PROCESS_COUNT 4 // Number of processes

// Function to perform Best Fit allocation
void bestFit(int partitions[], int partitionCount, int processes[], int
processCount) {
    int allocation[processCount];

    // Initially, no process is allocated
    for (int i = 0; i < processCount; i++) {
        allocation[i] = -1;
    }

    // Try to allocate memory for each process
    for (int i = 0; i < processCount; i++) {
        int minIndex = -1;

```

```

int minSize = 9999999;

// Find the partition with the smallest size that can accommodate the process
for (int j = 0; j < partitionCount; j++) {
    if (partitions[j] >= processes[i] && partitions[j] < minSize) {
        minSize = partitions[j];
        minIndex = j;
    }
}

// If a suitable partition is found, allocate it
if (minIndex != -1) {
    allocation[i] = minIndex;
    partitions[minIndex] -= processes[i]; // Allocate the partition
}

// Output the result
printf("\nBest Fit Allocation:\n");
for (int i = 0; i < processCount; i++) {
    if (allocation[i] != -1) {
        printf("Process %d allocated to Partition %d\n", i + 1, allocation[i] + 1);
    } else {
        printf("Process %d not allocated\n", i + 1);
    }
}

int main() {

    int partitions[PARTITION_COUNT] = {100, 500, 200, 300, 600}; // Partition
    sizes

```

```

int processes[PROCESS_COUNT] = {212, 417, 112, 426};    // Process sizes

bestFit(partitions, PARTITION_COUNT, processes, PROCESS_COUNT);

return 0;

}

```

**Experiment 9: Simulate the following page replacement algorithms a) FIFO  
b) LRU c) LFU**

### **1. FIFO (First-In-First-Out)**

```

#include <stdio.h>

#include <stdlib.h>

#define FRAME_COUNT 4 // Number of frames in memory

// Function to simulate FIFO page replacement
void fifo(int pages[], int pageCount) {
    int frames[FRAME_COUNT];
    int pageFaults = 0;
    int i, j;
    int isPresent;

    // Initialize frames to -1 (empty)
    for (i = 0; i < FRAME_COUNT; i++) {
        frames[i] = -1;
    }

    printf("FIFO Page Replacement:\n");

    for (i = 0; i < pageCount; i++) {
        isPresent = 0;

```

```

        // Check if the page is already in memory
        for (j = 0; j < FRAME_COUNT; j++) {
            if (frames[j] == pages[i]) {
isPresent = 1;
                break;
            }
        }

        // Page fault if the page is not in memory
        if (!isPresent) {
frames[pageFaults % FRAME_COUNT] = pages[i]; // Replace the oldest
page
pageFaults++;
        }

        // Print current memory state
printf("After accessing page %d: ", pages[i]);
        for (j = 0; j < FRAME_COUNT; j++) {
            if (frames[j] == -1) {
printf(" _ ");
            } else {
printf(" %d ", frames[j]);
            }
        }
        printf("\n");
    }

```

```
printf("\nTotal Page Faults: %d\n", pageFaults);  
}
```

```
int main() {  
    int pages[] = {7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 3};  
    int pageCount = sizeof(pages) / sizeof(pages[0]);  
  
    fifo(pages, pageCount);  
    return 0;  
}
```

## **2. LRU (Least Recently Used)**

```
#include <stdio.h>  
#include <stdlib.h>  
  
#define FRAME_COUNT 4 // Number of frames in memory  
  
// Function to simulate LRU page replacement  
void lru(int pages[], int pageCount) {  
    int frames[FRAME_COUNT];  
    int frequency[FRAME_COUNT] = {0}; // Frequency of page access  
    int pageFaults = 0;  
    int i, j;  
    int isPresent, lruIndex, minFrequency;  
  
    // Initialize frames to -1 (empty)  
    for (i = 0; i < FRAME_COUNT; i++) {  
        frames[i] = -1;  
    }  
}
```

```

printf("LFU Page Replacement:\n");

    for (i = 0; i < pageCount; i++) {
isPresent = 0;

        // Check if the page is already in memory
        for (j = 0; j < FRAME_COUNT; j++) {
            if (frames[j] == pages[i]) {
isPresent = 1;

                frequency[j]++; // Increment frequency for this page
                break;
            }
        }

        // Page fault if the page is not in memory
        if (!isPresent) {
            // Find the least frequently used page
            lfuIndex = 0;
            minFrequency = frequency[0];

            for (j = 1; j < FRAME_COUNT; j++) {
                if (frequency[j] < minFrequency) {
                    minFrequency = frequency[j];
                    lfuIndex = j;
                }
            }

            frames[lfuIndex] = pages[i];
        }
    }
}

```

```

        frequency[lfuIndex] = 1; // Set frequency to 1 for the new page
    pageFaults++;
}

```

```

    // Print current memory state
    printf("After accessing page %d: ", pages[i]);
    for (j = 0; j < FRAME_COUNT; j++) {
        if (frames[j] == -1) {
            printf(" _ ");
        } else {
            printf(" %d ", frames[j]);
        }
    }
    printf("\n");
}

```

```

printf("\nTotal Page Faults: %d\n", pageFaults);
}

```

```

int main() {
    int pages[] = {7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 3};
    int pageCount = sizeof(pages) / sizeof(pages[0]); // Missing semicolon
    here

```

```

    lfu(pages, pageCount);
    return 0;
}

```

### 3. LFU (Least Frequently Used)



```

#include <stdio.h>

#include <stdlib.h>

#define FRAME_COUNT 4 // Number of frames in memory

// Function to simulate LFU page replacement
void lfu(int pages[], int pageCount) {
    int frames[FRAME_COUNT];
    int frequency[FRAME_COUNT] = {0}; // Frequency of page access
    int pageFaults = 0;
    int i, j;
    int isPresent, lfuIndex, minFrequency; // Added semicolon here

    // Initialize frames to -1 (empty)
    for (i = 0; i < FRAME_COUNT; i++) {
        frames[i] = -1;
    }

    printf("LFU Page Replacement:\n");

    for (i = 0; i < pageCount; i++) {
        isPresent = 0;

        // Check if the page is already in memory
        for (j = 0; j < FRAME_COUNT; j++) {
            if (frames[j] == pages[i]) {
                isPresent = 1;
            }
        }
    }
}

```

```

        frequency[j]++; // Increment frequency for this page
        break;
    }
}

// Page fault if the page is not in memory
if (!isPresent) {
    // Find the least frequently used page
    lfuIndex = 0;
    minFrequency = frequency[0];
    for (j = 1; j < FRAME_COUNT; j++) {
        if (frequency[j] < minFrequency) {
            minFrequency = frequency[j];
            lfuIndex = j;
        }
    }
    frames[lfuIndex] = pages[i];
    frequency[lfuIndex] = 1; // Set frequency to 1 for the new page
    pageFaults++;
}

// Print current memory state
printf("After accessing page %d: ", pages[i]);
for (j = 0; j < FRAME_COUNT; j++) {
    if (frames[j] == -1) {
        printf(" _ ");
    } else {

```

```
printf(" %d ", frames[j]);
```

```
    }
```

```
    }
```

```
printf("\n");
```

```
}
```

```
printf("\nTotal Page Faults: %d\n", pageFaults);
```

```
}
```

```
int main() {
```

```
    int pages[] = {7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 3};
```

```
    int pageCount = sizeof(pages) / sizeof(pages[0]);
```

```
    lfu(pages, pageCount);
```

```
    return 0;
```

```
}
```

### **Experiment 10: Simulate Paging Technique of memory management.**

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
#define FRAME_COUNT 4 // Number of frames in physical memory
```

```
#define PAGE_COUNT 8 // Number of pages in the process
```

```
#define PAGE_SIZE 4 // Size of each page (number of addresses per  
page)
```

```
// Page Table: Maps page numbers to frame numbers
```

```
int pageTable[PAGE_COUNT] = {-1}; // Initialize all page entries as -1 (not  
allocated)
```

```

// Function to simulate page allocation and memory mapping
void pagingSimulation(int processPages[], int pageCount) {
    int physicalMemory[FRAME_COUNT][PAGE_SIZE] = {0}; // Physical
memory divided into frames

    int pageFaults = 0;

    int i, pageIndex, frameIndex;

    printf("Paging Simulation:\n");

    for (i = 0; i < pageCount; i++) {
        pageIndex = processPages[i];

        // Check if the page is already mapped to a frame (no page fault)
        if (pageTable[pageIndex] != -1) {
            frameIndex = pageTable[pageIndex];
            printf("Page %d is already in frame %d\n", pageIndex, frameIndex);
        } else {
            // Page fault, find an empty frame
            int foundEmptyFrame = 0;
            for (frameIndex = 0; frameIndex < FRAME_COUNT; frameIndex++)
            {
                if (physicalMemory[frameIndex][0] == 0) { // Frame is empty
                    // Allocate the page to the frame
                    pageTable[pageIndex] = frameIndex;
                    physicalMemory[frameIndex][0] = pageIndex;
                    foundEmptyFrame = 1;
                }
            }
            pageFaults++;
        }
    }
}

```

```
printf("Page %d caused a page fault. Allocating to frame %d\n", pageIndex,
frameIndex);
```

```
    break;
```

```
    }
```

```
}
```

```
    if (!foundEmptyFrame) {
```

```
        // If no empty frame, replace the first frame (simple replacement
policy)
```

```
frameIndex = 0;
```

```
pageTable[pageIndex] = frameIndex;
```

```
physicalMemory[frameIndex][0] = pageIndex;
```

```
printf("Page %d caused a page fault. Replacing frame %d\n", pageIndex,
frameIndex);
```

```
pageFaults++;
```

```
    }
```

```
}
```

```
    // Display the page table and physical memory status
```

```
printf("Current Page Table:\n");
```

```
    for (int j = 0; j < PAGE_COUNT; j++) {
```

```
printf("Page %d -> Frame %d\n", j, pageTable[j]);
```

```
    }
```

```
printf("Current Physical Memory:\n");
```

```
    for (int j = 0; j < FRAME_COUNT; j++) {
```

```
        if (physicalMemory[j][0] != 0) {
```

```
printf("Frame %d: Page %d\n", j, physicalMemory[j][0]);
```

```
        } else {
```

```
printf("Frame %d: Empty\n", j);
```

```
}
```

```
}
```

```
printf("\n");
```

```
}
```

```
printf("Total Page Faults: %d\n", pageFaults);
```

```
}
```

```
int main() {
```

```
    // Simulate a process with page references (logical addresses)
```

```
    int processPages[] = {2, 4, 1, 2, 5, 6, 3, 2, 4, 1}; // Logical page accesses
```

```
    int pageCount = sizeof(processPages) / sizeof(processPages[0]); // Added  
    semicolon here
```

```
    pagingSimulation(processPages, pageCount); // Added semicolon here
```

```
    return 0;
```

```
}
```

### **Experiment 11:Implement Bankers Algoriocat Dest Lock avoidance and prevention**

```
#include <stdio.h>
```

```
#include <stdbool.h>
```

```
#define P 5 // Number of processes
```

```
#define R 3 // Number of resources
```

```

// Function to check if a state is safe using Banker's Algorithm

bool isSafeState(int processes[], int avail[], int max[][R], int allot[][R], int
need[][R]) {

    int work[R], finish[P];

    int safeSeq[P];

    // Initialize work[] and finish[] arrays
    for (int i = 0; i < R; i++) {
        work[i] = avail[i];
    }
    for (int i = 0; i < P; i++) {
        finish[i] = 0;
    }

    int count = 0; // Number of processes that can finish
    while (count < P) {
        bool found = false;
        for (int p = 0; p < P; p++) {
            if (finish[p] == 0) {
                bool canFinish = true;

                // Check if the process can be finished (all needed resources are
                available)
                for (int j = 0; j < R; j++) {
                    if (need[p][j] > work[j]) {
                        canFinish = false;
                        break;
                    }
                }
            }
        }
    }
}

```

```

        if (canFinish) {
            // Add allocated resources to work[] and mark the process as
finished
            for (int j = 0; j < R; j++) {
                work[j] += allot[p][j];
            }
            finish[p] = 1; // Mark process p as finished
            safeSeq[count++] = p;
            found = true;
        }
    }
    if (!found) {
        return false; // No process can finish, system is in unsafe state
    }
}

// Print the safe sequence
printf("Safe Sequence: ");
    for (int i = 0; i < P; i++) {
        printf("P%d ", safeSeq[i]);
    }
    printf("\n");

    return true; // System is in a safe state
}

```

// Function to request resources using Banker's Algorithm



```

bool requestResources(int p, int request[], int processes[], int avail[], int
max[][R], int allot[][R], int need[][R]) {
    // Check if the request is less than or equal to the needed resources
    for (int i = 0; i < R; i++) {
        if (request[i] > need[p][i]) {
printf("Error: Process has exceeded its maximum claim\n");
            return false;
        }
    }
    // Check if the request is less than or equal to the available resources
    for (int i = 0; i < R; i++) {
        if (request[i] > avail[i]) {
printf("Error: Resources are not available\n");
            return false;
        }
    }
    // Pretend to allocate resources and check if the system is still in a safe
state
    for (int i = 0; i < R; i++) {
        avail[i] -= request[i];
        allot[p][i] += request[i];
        need[p][i] -= request[i];
    }

    // Now check if the system is still in a safe state after allocation
    if (isSafeState(processes, avail, max, allot, need)) {
printf("Resources allocated successfully\n");
        return true;
    }
}

```

```

    } else {
        // Rollback the allocation if not in a safe state
        for (int i = 0; i < R; i++) {
            avail[i] += request[i];
            allot[p][i] -= request[i];
            need[p][i] += request[i];
        }
        printf("Resources could not be allocated due to unsafe state\n");
        return false;
    }
}

```

```

int main() {
    int processes[] = {0, 1, 2, 3, 4}; // Process IDs
    int avail[] = {3, 3, 2}; // Available resources

    // Maximum resources needed by each process
    int max[P][R] = {
        {7, 5, 3},
        {3, 2, 2},
        {9, 0, 2},
        {2, 2, 2},
        {4, 3, 3}
    };

    // Resources allocated to each process
    int allot[P][R] = {

```

```
    {0, 1, 0},  
    {2, 0, 0},  
    {3, 0, 2},  
    {2, 1, 1},  
    {0, 0, 2}  
};
```

```
// Calculate the need matrix (max - allot)
```

```
int need[P][R];  
for (int i = 0; i < P; i++) {  
    for (int j = 0; j < R; j++) {  
        need[i][j] = max[i][j] - allot[i][j];  
    }  
}
```

```
// Check if the system is in a safe state initially
```

```
if (isSafeState(processes, avail, max, allot, need)) {  
printf("System is in a safe state\n");  
    } else {  
printf("System is in an unsafe state\n");  
    }
```

```
// Request resources for process 1 (e.g., request 1 unit of resource 0, 0 of  
resource 1, 2 of resource 2)
```

```
int request[] = {1, 0, 2};  
requestResources(1, request, processes, avail, max, allot, need);
```

```
return 0;
```

```
}
```

**Experiment 12: Simulate the following file allocation strategies a) Sequential b) Indexed e) Linked explain all with code**

**a) Sequential Allocation:**

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
#define MAX_BLOCKS 10 // Maximum number of blocks in the file system
```

```
typedef struct {
```

```
    int block[MAX_BLOCKS]; // Represents the storage blocks
```

```
    int allocated[MAX_BLOCKS]; // Allocation status (1 for allocated, 0 for free)
```

```
} Disk;
```

```
void sequentialAllocation(Disk *disk, int fileSize) {
```

```
    int blocksNeeded = (fileSize + MAX_BLOCKS - 1) / MAX_BLOCKS;
```

```
    int i;
```

```
    for (i = 0; i < blocksNeeded; i++) {
```

```
        if (disk->allocated[i] == 1) {
```

```
printf("Block %d already allocated!\n", i);
```

```
        } else {
```

```
            disk->allocated[i] = 1;
```

```
printf("Allocating block %d to file\n", i);
```

```
        }
```

```
    }
```

```
} // <-- Closing brace for sequentialAllocation function
```

```

int main() {
    Disk disk = {0};

    int fileSize = 15; // Size of the file to be allocated in blocks
    sequentialAllocation(&disk, fileSize);

    return 0;
}

```

### **b) Indexed Allocation:**

```

#include <stdio.h>

#include <stdlib.h> // Closing bracket added here.

#define MAX_BLOCKS 10 // Maximum number of blocks in the file
system

#define INDEX_SIZE 3 // Size of the index block

typedef struct {
    int block[MAX_BLOCKS]; // Represents the storage blocks

    int allocated[MAX_BLOCKS]; // Allocation status (1 for allocated, 0 for
free)

    int index[INDEX_SIZE]; // Index block with pointers to data blocks
} Disk;

void indexedAllocation(Disk *disk, int fileSize) {
    int blocksNeeded = (fileSize + MAX_BLOCKS - 1) / MAX_BLOCKS;
    int i;
    for (i = 0; i < blocksNeeded; i++) {
        if (disk->allocated[i] == 1) {
            printf("Block %d already allocated!\n", i);

```

```

    } else {
        disk->allocated[i] = 1;
        disk->index[i % INDEX_SIZE] = i; // Store block index in the index
        block
        printf("Allocating block %d to file (Index: %d)\n", i, i % INDEX_SIZE);
    }
}
}

```

```

int main() {
    Disk disk = {0};
    int fileSize = 12; // Size of the file to be allocated in blocks
    indexedAllocation(&disk, fileSize);
    return 0;
}

```

### c) **Linked Allocation:**

```

#include <stdio.h>
#include <stdlib.h>

#define MAX_BLOCKS 10 // Maximum number of blocks in the file
system

typedef struct {
    int block[MAX_BLOCKS]; // Represents the storage blocks
    int allocated[MAX_BLOCKS]; // Allocation status (1 for allocated, 0 for
    free)
    int next[MAX_BLOCKS]; // Pointers to the next block in the chain
} Disk;

void linkedAllocation(Disk *disk, int fileSize) {
    int blocksNeeded = (fileSize + MAX_BLOCKS - 1) / MAX_BLOCKS;

```

```

int i, prevBlock = -1;
for (i = 0; i < blocksNeeded; i++) {
    int blockIndex = -1;
    for (int j = 0; j < MAX_BLOCKS; j++) {
        if (disk->allocated[j] == 0) {
            blockIndex = j;
            break;
        }
    }
    if (blockIndex == -1) {
        printf("No available blocks for allocation\n");
        return;
    } else {
        disk->allocated[blockIndex] = 1;
        if (prevBlock != -1) {
            disk->next[prevBlock] = blockIndex; // Set pointer from previous block
        }
        disk->next[blockIndex] = -1; // Last block points to -1 (end of chain)

        if (prevBlock == -1) {
            printf("Allocating block %d to file (Start of chain)\n", blockIndex);
        } else {
            printf("Allocating block %d to file (Linked to block %d)\n", blockIndex,
                prevBlock);
        }
        prevBlock = blockIndex;
    }
}

```

```
}  
  
int main() {  
    Disk disk = {0};  
    int fileSize = 6; // Size of the file to be allocated in blocks  
    linkedAllocation(&disk, fileSize);  
    return 0;  
}
```